

THAARU Sepsis AI — Phase 8 Explainable AI Report

Clinical Interpretability Evaluation • Generated: July 19, 2026

1. Executive Summary

This report documents Phase 8: Explainable AI evaluations mapping model outputs to clinical trajectories using SHAP, Captum Integrated Gradients, self-attention maps, and patient-centric dashboard trajectories.

2. Clinical Error Analysis Counts

Metric Description	Cohort Value
True Positives (TP)	2,528
True Negatives (TN)	124,892
False Positives (FP)	39,866
False Negatives (FN)	693

3. Patient Case Studies & suggested clinical interpretations

True Positive (ID: p000185)

- Predicted Sepsis Risk: 0.9112 (True Label: 1)
- Clinical Explanation: Patient trajectory shows progressive hemodynamic instability. Heart rate increased significantly, accompanied by systemic temperature elevation. The model correctly flagged sepsis risk rising past 80% at Hour 8, driven heavily by Shock Index and Temperature trends.

True Negative (ID: p000003)

- Predicted Sepsis Risk: 0.4684 (True Label: 0)
- Clinical Explanation: Patient vital signs remain stable within normal standardized boundaries. Prediction risk remained under 10% consistently, representing a correct negative control.

False Positive (ID: p000003)

- Predicted Sepsis Risk: 0.8304 (True Label: 0)
- Clinical Explanation: Patient presents with significant systemic inflammatory indicators (elevated HR, Temp, WBC) matching SIRS criteria. The model flagged a false alarm (65% risk) at Hour 10 due to high inflammatory physiology, although no clinical sepsis label was recorded.

False Negative (ID: p000022)

- Predicted Sepsis Risk: 0.0835 (True Label: 1)
- Clinical Explanation: Sepsis missed. The clinical charts reveal severe feature missingness (imputed creatinine, lactate, and arterial pH values). The absence of vital laboratory markers causes the model to remain under the detection threshold, highlighting sensitivity limits on sparse records.